


AI AGENT EVALS

74-78%

SWE-bench Verified capability score

A high benchmark score is not production readiness.



So what is the benchmark missing?

Sources: SWE-bench Verified; aggregated production PR-acceptance reports (2026).

THE PATH

From benchmark to release gate

Four checks close the behavior gap.

-
- ```
graph TD; A[1 Define the behavior contract] --> B[2 Build a real-task eval set]; B --> C[3 Grade behavior, not just answers]; C --> D[4 Gate every release in CI];
```
- 1 Define the behavior contract**
  - 2 Build a real-task eval set**
  - 3 Grade behavior, not just answers**
  - 4 Gate every release in CI**

Sources: SWE-bench Verified; aggregated production PR-acceptance reports (2026).

THE GAP IS BEHAVIORAL

## Can-solve is not will-follow

**74-78%**

SWE-bench Verified — can  
the agent solve the task

**35-50%**

Real PR acceptance — does  
it follow the contract


Benchmarks ask whether an agent can solve a task. Production evals ask whether it honors the behavior contract every time.

Sources: SWE-bench Verified; aggregated production PR-acceptance reports (2026).



## THE SYSTEM

# Four layers make evals real



Start small: one task suite, behavior graders, a CI gate, and history tracked by model, prompt, and release.

Source: first-party AI agent evals implementation pattern.

## STEP 1

# Use real workflows first

**38**

production workflow tasks

**8**

agents under test

Score the agent on the real tasks users run — not just curated benchmark prompts that flatter the model.

Source: first-party implementation metrics.

## STEP 2

# Grade the behavior trace

**[x]** Assert the expected tool calls fired

**[x]** Check refusal boundaries held

**[x]** Require confirmation gates before writes

Final-answer checks are not enough. A polished answer can hide a skipped step.

Source: OpenAI Evals guide; first-party behavior graders.

### STEP 3

## Block regressions in CI

**\$3-8**

per eval run

**200-  
400K**

tokens per run

**15-25m**

parallel runtime

A full behavior eval suite is cheap enough to run on every pull request — treat a failure like a failing test.

Source: first-party implementation metrics.



## PATTERN INTERRUPT

**The answer can look perfect while the agent quietly skipped the work.**

That is exactly the failure benchmarks never catch — and the one your users feel first.

Source: AI agent evals failure taxonomy.



## RECAP

# Your release control panel

[x]

Real-task eval set

[x]

Behavior + tool-call graders

[x]

CI gate wired to block ships

[ ]

Ship gate missing — scores are decoration

Source: AI agent evals production-readiness guide.

TAKE THIS WITH YOU

## Save this before your next agent release

The Sourdough Test turns silent persona drift into a visible, gated release signal.

**Guide** Full visual walkthrough — link in comments.

Source: AI agent evals production-readiness series.